



Innovative. Technology. Partner.

Sharp Patient Risk Score POC – Technology Transfer Workshop



Overview

- Housekeeping
- Review the project objectives
- Discuss the methodologies
- Discuss the result
- Demonstrate and discuss the project artifacts
 - Impala/Hive scripts for querying the Cerner tables in HDFS
 - Jupyter notebooks
 - Exploratory Data Analysis (EDA)
 - Data cleaning
 - Predictive model training and validation

Housekeeping

- You should have Anaconda (Python 2.7 version) already installed
- <https://www.continuum.io/downloads>
- The tech transfer directory and files are best viewed through a Jupyter notebook, which is easily obtained by installing Anaconda
- **Does anyone not have access to either the GitHub repository or the tech-transfer files?**

Project Review

- **Context** – Sharp HealthCare is a not-for-profit integrated regional health care delivery system in San Diego that consists of:
 - four acute-care hospitals
 - three specialty hospitals
 - two affiliated medical groups
 - a health plan
- **Goal** - to build a Proof of Concept implementation capable of identifying patients at risk for an adverse rapid response team (RRT) event

Project Flow

- Project Participants:
 - Intel (project management and coordination)
 - ProKarma (analytics and development)
 - Sharp Healthcare (subject matter expertise and data owner)
- Initial Week consisted of a daily meeting during discovery and setup phase
- Subsequent weeks (Aug 19th – Sep 30th) consisted of weekly sprint reviews with Intel, PK, and Sharp personnel
- **Final Deliverables:**
 - **Knowledge transfer workshop (now)**
 - **Present project findings to Sharp Executive team**

Methodologies

- Identify which patients had RRT events
- Identify patient features that are plausible causal factors for a sudden decline in health status
 - E.g. vital signs, medication usage, narcotics, movements between units, etc.
- Extract the subsets of data from which the desired patient features can be derived
- Process and analyze the features, while identifying which are sufficiently dense for model training
 - Select a cohort of counter-examples, i.e. Who did not have an RRT event?
- Model training and selection via cross validation

Results

- We were able to create probabilistic risk scores for patients
- Against validation dataset:
 - Accuracy = 80% (number correctly classified to their respective class)
 - Precision = 82% (positive predictive value)
- Performance improved by
 - Adding more training data, 0.5 years to 1.5 years of data
 - Adding more varied features
 - Further model optimization
- Performance impeded by
 - Lack of training instances
 - Lack of dense data features

Predictive signal definitely exists and our work suggests a clear path to improve performance

Project Artifacts

- The workflow
 - Use Impala to query Cerner data in Hive format
 - Conduct analysis and modeling using Anaconda in Jupyter notebooks
 - Mostly using Pandas and scikit-learn
- The structure of the project files
 - OVERVIEW.ipynb
 - etl-queries (directory)
 - notebooks (directory)
 - EDA (directory)
 - modeling (directory)

Demonstrate and Discuss

- Let's see some code!



Questions